

Topic: MPI Programming

Objective

- MPI point-to-point communication
- Important points to remember while implementing MPI
- Hands-on

Topic: MPI Programming

Blocking and non-blocking


Blocking	Non-blocking
MPI_Send	MPI_Isend
MPI_Recv	MPI_Irecv

Blocking: the process does not return until data transmitted is started from the buffer

Non-blocking: the process return immediately after the op

Topic: MPI Programming

Sending data in a ring-like pattern/topology



Topic: MPI Programming

MPI_Isend/MPI_Irecv

```
program test
  implicit none
  include 'mpif.h'

  integer :: p, id, err, root, msg, tag, request(MPI_Status_Size)

  call MPI_Init(err)
  call MPI_Comm_Size(MPI_Comm_World, p, err)
  call MPI_Comm_Rank(MPI_Comm_World, id, err)

  root=0 ; tag=0
  if(id==root) then
    msg=10
    call MPI_Isend(msg,1,MPI_Int,1,tag,MPI_Comm_World,request,err)
  else
    call MPI_Irecv(msg,1,MPI_Int,id-1,MPI_Any_Tag,MPI_Comm_World,request,err)
    write(*,*) id,'received from process:',id-1,'msg: ',msg
    call MPI_Isend(msg,1,MPI_Int,mod(id+1,p),tag,MPI_Comm_World,request,err)
  endif

  if(id==root) then
    call MPI_Irecv(msg,1,MPI_Int,p-1,MPI_Any_Tag,MPI_Comm_World,request,err)
    write(*,*) id,'received from process:',p-1,'msg: ',msg
  endif

  call MPI_Finalize(err)
end program test
```

Topic: MPI Programming

Output

```
mpirun -np 4 ./mpifring.x
0 received from process:      3 msg:      10
2 received from process:      1 msg:       0
1 received from process:      0 msg:       0
3 received from process:      2 msg:       0
```

Topic: MPI Programming

Correct – MPI_Isend/MPI_Irecv

```
program test
  implicit none
  include 'mpif.h'

  integer :: p, id, err, root, msg, tag, request(MPI_Status_Size), status(MPI_Status_Size), msg1

  call MPI_Init(err)
  call MPI_Comm_Size(MPI_Comm_World, p, err)
  call MPI_Comm_Rank(MPI_Comm_World, id, err)

  root=0 ; tag=0
  if(id==root) then
    msg=10
    call MPI_Isend(msg,1,MPI_Int,1,tag,MPI_Comm_World,request,err)
  else
    call MPI_Irecv(msg1,1,MPI_Int,id-1,MPI_Any_Tag,MPI_Comm_World,request,err)
    call MPI_Wait(request,status)
    write(*,*) id,'received from process:',id-1,'msg: ',msg1
    call MPI_Isend(msg1,1,MPI_Int,mod(id+1,p),tag,MPI_Comm_World,request,err)
  endif

  if(id==root) then
    call MPI_Irecv(msg1,1,MPI_Int,p-1,MPI_Any_Tag,MPI_Comm_World,request,err)
    call MPI_Wait(request,status)
    write(*,*) id,'received from process:',p-1,'msg: ',msg1
  endif

  call MPI_Finalize(err)
end program test
```

Topic: MPI Programming

Can we parallelize all parts of the programs?

- Fibonacci series: Generate 'N' Fibonacci numbers and calculate their average (can we implement MPI, here?)

(

Sum of previous two numbers gives the number

$$x_i = x_{(i-1)} + x_{(i-2)}$$

)

Topic: MPI Programming

Can we parallelize all part of the programs? **NO**

!Fibonacci series: each number will be the sum of the two preceding numbers

```
program fibo_series
  implicit none

  integer(kind=16) :: sum,i,first,second,third,N

  write(*,*) 'enter N value'
  read(*,*) N

  first=1 ; second=1

  do i=1,N
    third=first+second
    sum=sum+third

    first=second
    second=third
    write(*,*) third
  enddo

end program fibo_series
```

When the current process depends on the data in other processes, synchronization needs to be established among processes, hence the communication time may dominate.

Topic: MPI Programming

Hands on

- Using the above data 'random_numbers.dat', write a program to calculate the average value of N numbers. a) Implement MPI using point-to-point blocking communication protocols b) Implement MPI using point-to-point non-blocking communication protocols. Show that the result in both cases is the same.